

## Exercise 1

- The modified part of the [code](#) along with the [output](#) is shown [below](#)

```

9 // in-line assembly to increment the x by 1
10 __asm (
11     "incl %0 \n"
12     : "=r" (x) // Output registers
13     : "0" (x) // Input registers
14 );

```

```

sking@sking:~/Documents/FIFTH SEMESTER IITG/OS Lab/Assgn_0/Submission$ gcc ex1.c
sking@sking:~/Documents/FIFTH SEMESTER IITG/OS Lab/Assgn_0/Submission$ ./a.out
Hello x = 1
Hello x = 2 after increment
OK
sking@sking:~/Documents/FIFTH SEMESTER IITG/OS Lab/Assgn_0/Submission$

```

Fig 1.1: in-line assembly to increment x

- `__asm__` function takes in **three** arguments separated by ':' and are described below
  - `incl %0 \n`: **Increases** the value of register `%0` by **1** (1 indicates **32 bit** register)
  - `"=r" (x)`: **Dynamically** allocate an **output** register for variable `x` (referred by `%0`)
  - `"0" (x)`: Use the register `%0` as an **input** register for variable `x`
- From the [second line of output](#) it is clear that the value of `x` has been **incremented** (refer [ex1.c](#) code file for details).

## Exercise 2 (pre - tasks)

- After [cloning QEMU](#) and installing the dependencies following [output](#) is obtained as a result of `make qemu` command followed by `ls` command on the [VGA](#):

```

sking@sking:~/Documents/FIFTH SEMESTER IITG/OS Lab/xv6-public$ make qemu
qemu-system-i386 -serial mon:stdio -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512
i386...
cpu0: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 minodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
ls
1 1 512
.. 1 1 512
README 2 2 2288
cat 2 3 16280
echo 2 4 15136
forktest 2 5 9440
grep 2 6 18500
init 2 7 15720
kill 2 8 15164
ln 2 9 15016
ls 2 10 17648
mkdir 2 11 15264
rm 2 12 15240
sh 2 13 27876
stressfs 2 14 16152
userstests 2 15 67200
wc 2 16 17016
zombie 2 17 14832
console 3 18 0

```

Fig 2.1: Launching QEMU

- The [background](#) window is the [ubuntu terminal](#) and the [smaller](#) window is the [QEMU VGA](#). The output of `ls` command appears on **both** the screens which demonstrates that both the windows' **input** and **output** are **synced**.

## Exercise 2

- The [screenshot of final output](#) is attached below after which there is a detailed [explanation](#) of output and also of [each step](#) which was [followed](#) to get the output. (The labels at the left and right correspond to the serial no. used for explaining)

```

1 { sking@sking:~/Documents/FIFTH SEMESTER IITG/OS Lab/xv6-public$ make qemu-gdb
-gdb
*** Now run 'gdb'.
qemu-system-i386 -serial mon:stdio -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512 -S -gdb tcp:26000

(gdb) source .gdbinit
+ target remote localhost:26000
warning: No executable has been specified and target does not support
determining executable automatically. Try using the "file" command.
The target architecture is set to "i386".
[000:ffff] 0x7ff0: jmp 0x3630,0xf00e05b
breakpoint in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configuration of GDB. Attempting to continue with the default i386 settings.

(gdb) si
[000:e05b] 0x7ff0: cmpw 0xffc8,%cs(%esi)
breakpoint in ?? ()
(gdb) si
[000:e062] 0x7ff0: jne 0x7ff0e067
breakpoint in ?? ()
(gdb) si
[000:e066] 0x7ff0: xor %edx,%edx
breakpoint in ?? ()
(gdb) si
[000:e068] 0x7ff0: mov %edx,%ss
breakpoint in ?? ()
(gdb) si
[000:e06a] 0x7ff0: mov 0x7000,%sp
breakpoint in ?? ()
(gdb) si
[000:e070] 0x7ff0: mov 0x7c4,%dx
breakpoint in ?? ()
(gdb) si
[000:e076] 0x7ff0: jmp 0x7ff0e07a
breakpoint in ?? ()
(gdb) si
[000:cf24] 0x7ff0: cli
breakpoint in ?? ()
(gdb) si
[000:cf25] 0x7ff0: cld
breakpoint in ?? ()

```

Fig 2.2: Tracing first few instructions of the ROM BIOS

- Compare the operands
- Conditionally jump if result of previous instruction is not equal
- Set `%edx` to 0
- Set `%ss` to 0
- Set `%sp` to 0x7000
- Set `%dx` to 0x7c4
- Jump to the specified location.
- Clear the interrupt flag.
- Clear the direction flag.

- Command executed (in the left terminal window): `make qemu-gdb`  
Output: Blank QEMU VGA with a message - "Guest has not initialized the display yet"  
Explanation: This command is used to **debug** and **tweak** into **step by step** process for the **OS** to **boot** up. As a consequence, even the **VGA** has not been initialized yet.
- Command executed (in the right gdb terminal window): `source .gdbinit`  
Output: GDB session has started and **first** instruction is present at the location `[000:ffff]` and is a **jump** instruction.  
Explanation: The first command to be executed is at the **end of the BIOS** section of memory (precisely speaking **16 bits** from the end). This command causes a **jump to the first instruction of boot-up section of BIOS** (i.e., at the location `[000:e05b]`) from where BIOS tries to boot-up the **OS** into the system. This mechanism **allows** the BIOS to **re-execute** in case of any **failure** (as it towards the end of the BIOS section of memory).
- Command executed (in the right gdb terminal window): `si`  
Output: **Assembly instructions** of boot-up process in BIOS.

**Explanation:** BIOS is trying to set up the **Interrupt Descriptor Table** in the first few commands which is used to initialize devices like the VGA display. The instructions in the screenshot set the **stack segment register (%ss)** and the **stack pointer register (%sp)** with the appropriate values. The in detail meaning of steps are written at the **right** side of the image attached.

### Exercise 3

- The first part describes the setting up of the **breakpoint at 0x7c00** address location and tracing the source code in **bootasm.S** (left part of the image), disassembly in **bootblock.asm** (middle part of the image) and in the **GDB terminal** (right part of the image). Below the screenshot is a description of the tasks performed.

The screenshot shows three panels in a GDB interface:

- Left Panel (bootasm.S):** Source code for bootasm.S. Line 14 is circled in red with a red '4' and an arrow pointing to it. The code includes instructions like `.code16`, `.globl start`, `start:`, `cli`, and various `movw`, `movb`, `inb`, `outb` instructions.
- Middle Panel (bootblock.asm):** Disassembly of bootblock.asm. It shows instructions like `xorw %ax,%ax`, `7c01: 31 c0`, `movw %ax,%ds`, `7c03: 8e d8`, `movw %ax,%es`, `7c05: 8e c0`, `movw %ax,%ss`, `7c07: 8e d0`, and physical address lines like `00007c09 <seta20.1>`.
- Right Panel (GDB Terminal):** GDB commands and output. It shows `(gdb) b *0x7c00`, `Breakpoint 1 at 0x7c00`, `(gdb) c`, `Continuing.`, `[ 0:7c00] => 0x7c00: cli`, `Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()`, `(gdb) x/10i`, and a list of 10 disassembled instructions starting from `0x7c01: xor %eax,%eax` down to `0x7c13: in $0x64,%al`.

Red handwritten annotations on the right side of the GDB terminal output include a bracket grouping the last three instructions (lines 3, 2, 1 from bottom) and the numbers 1, 2, 3.

Fig 3.1:

- Setting breakpoint at 0x7c00.
- Comparing assembly, disassembly and output of x/10i.

- Command:** `b *0x700`  
**Explanation and Output:** This is used to set **break-point** at the instruction to be executed at memory location with address 0x7c00. The output tells that the breakpoint is **successfully** set.
  - Command:** `c`  
**Explanation and Output:** It **executes** all the instructions up to the next break-point (here at address 0x7c00). After the execution is done, it **prints the next** instruction to be executed onto the terminal (here `cli` instruction).
  - Command:** `x/10i`  
**Output:** **Ten lines** are printed and each line constitutes of a **disassembly** instruction.  
**Explanation:** This command **prints the next 10 disassembly instructions** to be executed. In general 10 can be replaced with **any positive integer**.
  - The **disassembly** code and **assembly** source code are **compared** with the disassembly **output** of Step 3. in GDB window. It is observed that all the instructions starting from `cli` all the way up to `inb $0x64 %al` are **exactly the same** except a few changes in the syntax (**majorly** the suffix of **b/w/l** standing for byte/word/long word is **not** present in GDB output).
- The next section of the exercise asks to locate various **parts of source code** and their corresponding **assembly** code. Following are the **screenshots** with labels for each

The screenshot shows two panels in a GDB interface:

- Left Panel (bootmain.c):** Source code for bootmain.c. It shows the `readsect` function with parameters `void *dst, uint offset`. The function includes comments like `// Issue command.` and `// Read data.` and instructions like `waitdisk()`, `outb(0x1F2, 1);`, `outb(0x1F3, offset);`, `outb(0x1F4, offset >> 8);`, `outb(0x1F5, offset >> 16);`, `outb(0x1F6, (offset >> 24) | 0xE0);`, `outb(0x1F7, 0x20);`, `insl(0x1F0, dst, SECTSIZE/4);`.
- Right Panel (bootblock.asm):** Disassembly of bootblock.asm. It shows instructions like `7c90: f3 0f 1e fb`, `7c94: 55`, `7c95: 89 e5`, `7c97: 57`, `7c98: 53`, `7c99: 8b 5d 0c`, `7c9c: e8 dd ff ff ff`, and `call 7c7e <waitdisk>`.

Fig 3.2:

- Left: Source code for readsect in bootmain.c
- Right: Corresponding assembly code in bootblock.asm

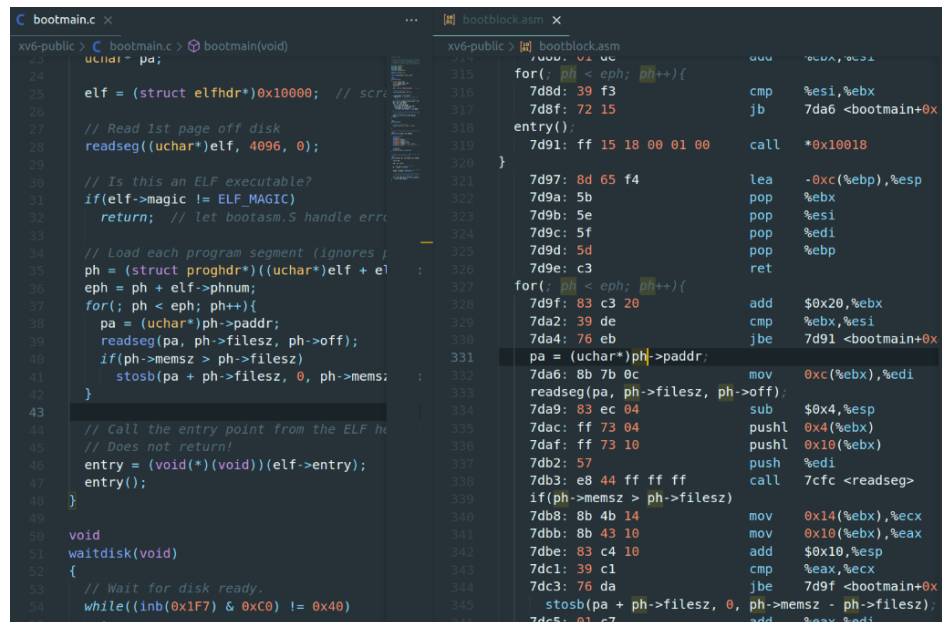


Fig.3.3: For loop reading remaining sectors

1. Left: Source code of for loop
2. Right: Assembly code of for loop

The code on line 46-47 (in left side source file of Fig 2.5) whose assembly is on line 318-319 (right side assembly file of Fig 2.5) of above image would be executed **once the for loop is over**. This line is the **final line** in execution of boot loader and after this the **control** is given to the OS. The **breakpoint** is set at 0x7d91 (line 319) as shown in Fig 2.6.

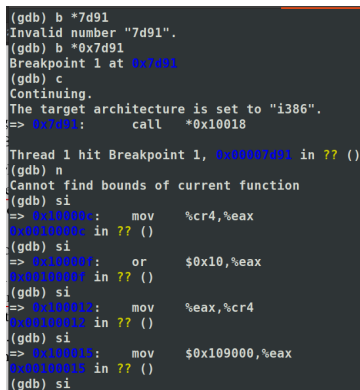


Fig.3.4: Breakpoint at 0x7d91 (end of for loop) and stepping through boot loader

- Finally the 3 questions are answered below:

1. Following image shows the last lines executed before switching to the protected mode. The **last line** to be executed is **line 51**. Before that the boot loader **deactivates the line A20** (to access 2MB of RAM instead of default 1MB) and also initializes the **descriptor table** which would be helpful in accessing the **virtual memory** in the **protected mode**. After it's done then the line 51 causes the control to **jump** to the protected mode.

Fig 3.5: Assembly code to switch from 16 bit Real mode to 32 bit Protected mode.

```
# Switch from real to protected mode. Use a bootstrap GDT that makes
# virtual addresses map directly to physical addresses so that the
# effective memory map doesn't change during the transition.
lgdt gdt_desc
movl %cr0, %eax
orl $CR0_PE, %eax
movl %eax, %cr0

//PAGEBREAK!
# Complete the transition to 32-bit protected mode by using a long jmp
# to reload %cs and %eip. The segment descriptors are set up with no
# translation, so that the mapping is still the identity mapping.
ljmp $(SEG_KCODE<<3), $start32
```

2. From the discussion in Fig 2.6 and from the source code at Fig 2.5, it is clear that the **last instruction** of boot loader would be to call the entry function (on line 47 in Fig 2.5 bootmain.c). Thus a **breakpoint** is set for the corresponding assembly instruction and the command **si** is used to **trace** the next instruction (first instruction of the kernel) and following conclusions are made based on the screenshot below:
  - **Last instruction of boot loader:** `call *0x10018`
  - **First instruction of kernel:** `mov %cr4, %eax`

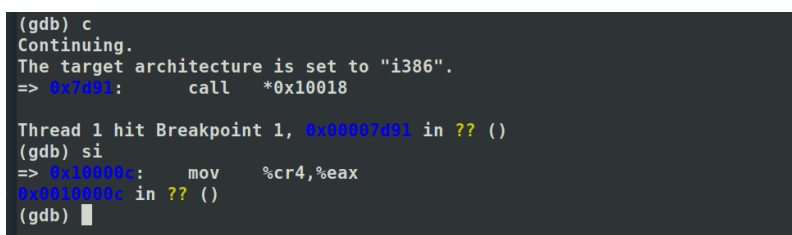


Fig 3.6: GDB terminal showing last instruction of boot loader and the first instruction of kernel

3. Fig 2.9 tells that the boot loader runs a **for loop starting at ph** and **ending at eph - 1**. The value of **ph** is obtained from the **ELF header** and the number of iterations are given by **elf->phnum**, thus **eph = ph + elf->phnum**.

Fig 3.7: For loop in boot loader which loads all the sectors

```
// Load each program segment (ignores ph flags).
ph = (struct proghdr*)((uchar*)elf + elf->phoff);
eph = ph + elf->phnum;
for(; ph < eph; ph++){
    pa = (uchar*)ph->paddr;
    readseg(pa, ph->filesz, ph->off);
    if(ph->memsz > ph->filesz)
        stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
}
```

## Exercise 4

- The **first part** requires to understand line by line the output of the file **pointers.c**. Therefore, below is the line by line explanation of each line of output.
  - Statement:** `printf("1: a = %p, b = %p, c = %p\n", a, b, c);`  
**Reason for output of line 1:** The variables **a** and **c** belong to **stack** memory address and thus store **close addresses**. **Malloc** is used to **dynamically** allocate address in the **heap** space and so the value stored in **b** is having a **totally different** address.

```
4 void f(void)
5 {
6     int a[4];
7     int *b = malloc(16);
8     int *c;
9     int i;
10
11     printf("1: a = %p, b = %p, c = %p\n", a, b, c);
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

```
sking@sking:~/Documents/FIFTH SEMESTER IITG/OS Lab/Assgn_0$ gcc pointers.c
sking@sking:~/Documents/FIFTH SEMESTER IITG/OS Lab/Assgn_0$ ./a.out
1: a = 0x7ffc80316d40, b = 0x559f095252a0, c = 0x7ffc80316d67
2: a[0] = 200, a[1] = 101, a[2] = 102, a[3] = 103
3: a[0] = 200, a[1] = 300, a[2] = 301, a[3] = 302
4: a[0] = 200, a[1] = 400, a[2] = 301, a[3] = 302
5: a[0] = 200, a[1] = 128144, a[2] = 256, a[3] = 302
6: a = 0x7ffc80316d40, b = 0x7ffc80316d44, c = 0x7ffc80316d41
sking@sking:~/Documents/FIFTH SEMESTER IITG/OS Lab/Assgn_0$
```

Fig 4.1: Output and source code for line 1

- The **statement** is attached as a **screenshot**.  
**Reason for output of line 2:** After line 13, the pointer **c** stores the **address of array a**. Thus line 15 changes the value of previously assigned **a[0]** of 100 to 200 and so in the output the value of **a[0]** is 200 whereas the **other values** of the array are **same** as before.

```
13     c = a;
14     for (i = 0; i < 4; i++)
15         a[i] = 100 + i;
16     c[0] = 200;
17     printf("2: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n",
18           a[0], a[1], a[2], a[3]);
```

Fig 4.2: Source code for line 2

- Reason for output of line 3:** Line 20 changes the value of **a[1]** to 300 (as **c** is still storing the address of array **a**). Lines 21 and 22 are just two other ways of **dereferencing** the elements of an array and change **a[2]** and **a[3]** to 301 and 302 respectively. Thus the only **difference** between output of line 2 and 3 is the values of **a[1]**, **a[2]** and **a[3]**.

```
20     c[1] = 300;
21     *(c + 2) = 301;
22     3 [c] = 302;
23     printf("3: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n",
24           a[0], a[1], a[2], a[3]);
```

Fig 4.3: Source code for line 3

- Reason for output of line 4:** Line 26 increases the address in **c** by 4 bytes (since **c** is of the type **int\***). Thus **c** now points to address of **a[1]** and Line 27 changes **a[1]** to 400. Hence the output in line 4 is same as that of line 2 except the value of **a[1]** is now 400.

```
26     c = c + 1;
27     *c = 400;
28     printf("4: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n",
29           a[0], a[1], a[2], a[3]);
```

Fig 4.4: Source code for line 4

- Reason for output of line 5:** Line 31 increases the address in **c** by 1 byte (since **c** is changed to **char\*** before incrementing). Now **c** doesn't point to any specific element in **a** and thus the value of 500 is assigned **partially between a[1] and a[2]** and it changes both the values in a **corrupted** way.

```
31     c = (int *)((char *)c + 1);
32     *c = 500;
33     printf("5: a[0] = %d, a[1] = %d, a[2] = %d, a[3] = %d\n",
34           a[0], a[1], a[2], a[3]);
```

Fig 4.5: Source code for line 5

6. Reason for output of line 6: Line 36 stores the address of **a[1]** in **b** and line 37 stores the address of **a** incremented by 1 byte (similar to line 31). Thus the value of **a** is same as output of line 1 and value of **b** is **a + 4 bytes** and value of **c** is **a + 1 byte**.

Fig 4.6: Source code for line 6

```
36      b = (int *)a + 1;
37      c = (int *)((char *)a + 1);
38      printf("6: a = %p, b = %p, c = %p\n", a, b, c);
```

- The next part asks us to solve the **K-splice** pointer challenge from which following **conclusions** are drawn:
  - Pointers and arrays are two **different** things, although **sometimes** arrays **decay** into pointers.
  - If **x** is an array of **int**, then the type of **&x** is a pointer to an array and **&x + 1** increments the address by **4\*n bytes** (where **n** is size of array **x**).
  - On the other hand **x + 1** increments the address by **4 bytes** (since in this case **x** decays into **&x[0]**).

Result: printf("%p\n", &x) prints 0x7ffdfbf7f100.

Aside: what is the type of &x[0]? Well, x[0] is an int, so &x[0] is "pointer to int". That feels right.

```
printf("%p\n", &x+1);
```

What will this print?  Guess

Ok, now for the coup de grace. **x** may be an array, but **&x** is definitely a pointer. So what's **&x+1**?

First, another aside: what is the type of **&x**? Well... **&x** is a pointer to an array of 5 ints. How would you declare something like that?

Let's fire up **cdecl** and find out:

```
cdecl> declare y as array 5 of int;
int y[5]
cdecl> declare y as pointer to array 5 of int;
int (*y)[5]
```

Confusing syntax, but it works:

```
int (*y)[5] = &x; compiles without error and works the way you'd expect.
```

But back to the question at hand. Pointer arithmetic tells us that **&x+1** is going to be the address of **x + sizeof(x)**. What's **sizeof(x)**? Well, it's an array of 5 ints. On this system, each int is 4 bytes, so it should be 20 bytes, or 0x14.

Result **&x+1** prints 0x7ffdfbf7f114.

Fig 4.7: K-Splice pointer challenge

## Exercise 5 (pre - tasks):

- Different **program sections** (especially **.text**, **.data** and **.rodata**) were analysed in the **ELF** binary file (**kernel.ld**). Screenshot of some of the sections is attached along with a brief explanation of observations:

```
.text : AT(0x100000) {
    *(.text .stub .text.* .gnu.linkonce.t.*)
}

PROVIDE(etext = .); /* Define the 'etext' symbol to this value */

.rodata : {
    *(.rodata .rodata.* .gnu.linkonce.r.*)
}
```

Fig 5.1: Important program sections of kernel.ld

- .text**: This section consists of program's **executable** instructions.
  - .data**: This section consists of program's **initialized data** like global variables.
  - .rodata**: This section consists of program's **read-only** data.
- The commands **objdump -h kernel** and **objdump -h bootblock.o** were executed to inspect various code sections in the **kernel** file and **text** section of the boot loader. The screenshot of the output is attached below.

```
skingsking:~/Documents/FIFTH SEMESTER IITG/OS Lab/xv6-public$ objdump -h kernel
kernel:      file format elf32-i386

Sections:
Idx Name          Size      VMA           LMA           File off  Algn
 0 .text          000070da  80100000  00100000  00001000  2**4
 1 .rodata         000009cb  801070e0  001070e0  000080e0  2**5
 2 .data           00002516  80108000  00108000  00009000  2**12
 3 .bss            0000af88  8010a520  0010a520  0000b516  2**5
 4 .debug_line     00006cb5  00000000  00000000  0000b516  2**0
 5 .debug_info     000121ce  00000000  00000000  000121cb  2**0
 6 .debug_abbrev   00003fd7  00000000  00000000  00024399  2**0
 7 .debug_ranges  000003a8  00000000  00000000  00028370  2**3
 8 .debug_str      00000ecf  00000000  00000000  00028718  2**0
 9 .debug_loc      0000681e  00000000  00000000  000295e7  2**0
10 .debug_ranges  00000d08  00000000  00000000  0002fe05  2**0
11 .comment        0000002a  00000000  00000000  00030b0d  2**0

skingsking:~/Documents/FIFTH SEMESTER IITG/OS Lab/xv6-public$ objdump -h bootblock.o
bootblock.o:  file format elf32-i386

Sections:
Idx Name          Size      VMA           LMA           File off  Algn
 0 .text          000001d3  00007c00  00007c00  00000074  2**2
 1 .eh_frame       000000b0  00007dd4  00007dd4  00000248  2**2
 2 .comment        0000002a  00000000  00000000  000002f8  2**0
 3 .debug_aranges  00000040  00000000  00000000  00000328  2**3
 4 .debug_info     00000542  00000000  00000000  00000368  2**0
 5 .debug_abbrev   0000022c  00000000  00000000  0000093a  2**0
 6 .debug_line     0000029a  00000000  00000000  00000b66  2**0
 7 .debug_str      00000244  00000000  00000000  00000e00  2**0
 8 .debug_loc      000002bb  00000000  00000000  00001044  2**0
 9 .debug_ranges  00000078  00000000  00000000  000012ff  2**0
```

Fig 5.2: objdump output for kernel (left) and bootblock.o (right)

## Exercise 5:

- On tracing through the first few instructions of boot loader, I concluded that the assembly instruction on **line 51** (**ljmp** instruction) would be the **first** one to **break** on **changing** the **link address** in the screenshot as shown below.

```

38
39  # Switch from real to protected mode. Use a bootstrap GDT that makes
40  # virtual addresses map directly to physical addresses so that the
41  # effective memory map doesn't change during the transition.
42  lgdt    gdt_desc
43  movl    %cr0, %eax
44  orl     $CR0_PE, %eax
45  movl    %eax, %cr0
46
47  //PAGEBREAK!
48  # Complete the transition to 32-bit protected mode by using a long jmp
49  # to reload %cs and %eip. The segment descriptors are set up with no
50  # translation, so that the mapping is still the identity mapping.
51  ljmp     $(SEG_KCODE<<3), $start32

```

Fig 5.3: Instruction first to break with wrong link address

- In the **line 107** the **correct** link address: **0x7c00** has been replaced with an **incorrect** link address: **0x7e00** as shown below in the screenshot. After that the commands **make clean** and **make** are run on the terminal.

```

104  bootblock: bootasm.S bootmain.c
105      $(CC) $(CFLAGS) -fno-pic -O -nostdinc -I. -c bootmain.c
106      $(CC) $(CFLAGS) -fno-pic -nostdinc -I. -c bootasm.S
107      $(LD) $(LDFLAGS) -N -e start -Ttext 0x7e00 -o bootblock.o bootasm.o bootmain.o
108      $(OBJDUMP) -S bootblock.o > bootblock.asm
109      $(OBJCOPY) -S -O binary -j .text bootblock.o bootblock
110      ./sign.pl bootblock

```

Fig 5.4: Changing the link address in the makefile

- Fig 5.5 compares the output of GDB terminal when **correct** link address was used with the case when **wrong** link address was used. From the screenshot, following conclusions can be drawn:
  - The instructions **before** the **ljmp** instruction are **same** in the boot-loader for both the cases.
  - The instructions **including** and **after** the **ljmp** instruction are **different**.
  - There is **no message** saying: "*The target architecture is set to be i386*" in the case with **wrong** link address (which happens after the **ljmp** instruction in the correct link address situation).

<pre> (gdb) si [ 0:7c25] =&gt; 0x7c25: or    \$0x1,%ax 0x00007c25 in ?? () (gdb) si [ 0:7c29] =&gt; 0x7c29: mov    %eax,%cr0 0x00007c29 in ?? () (gdb) si [ 0:7c2c] =&gt; 0x7c2c: ljmp   \$0xb866,\$0x87c31 0x00007c2c in ?? () (gdb) si The target architecture is set to "i386". =&gt; 0x7c31:    mov    \$0x10,%ax 0x00007c31 in ?? () (gdb) si =&gt; 0x7c35:    mov    %eax,%ds </pre>	<pre> (gdb) si [ 0:7c25] =&gt; 0x7c25: or    \$0x1,%ax 0x00007c25 in ?? () (gdb) si [ 0:7c29] =&gt; 0x7c29: mov    %eax,%cr0 0x00007c29 in ?? () (gdb) si [ 0:7c2c] =&gt; 0x7c2c: ljmp   \$0xb866,\$0x87e31 0x00007c2c in ?? () (gdb) si [f000:e05b] 0xfe05b: cmpw   \$0xffc8,%cs:(%esi) 0x0000e05b in ?? () (gdb) si [f000:e062] 0xfe062: jne    0xd241d0b2 0x0000e062 in ?? () </pre>
--	---

Fig 5.5: Differences between the correct link address (left) and wrong link address (right) cases.

After the experiment, the link address was corrected and the commands **make clean** and **make** were executed as instructed.

- Next, it is asked to **repeat** the commands of **objdump** in the terminal for kernel and for boot-loader. With reference to Fig 5.2 (new screenshot is not used here to avoid repetition), we can see that the **boot-loader** has the **VMA** (link address) and **LMA** (load address) exactly the **same** for various sections. However, in case of **kernel** the **link address** (expected execution point) is at a **higher memory address** (shifted by **0x8000000**) for each of the sections which are to be loaded. For ex: in case of .text program section, the LMA is 0x00100000 (i.e., **1MB**) whereas the VMA is 0x8010000.
- The **e\_entry** ELF header field is stores in **struct elfhdr.entry** as shown in the **left** side of the screenshot below (line 12). Also, the **entry** point address (0x0010000c) of the kernel as recorded by execution of the command **objdump -f kernel** is shown below.

<pre> 8  uchar elf[12]; 9  ushort type; 10 ushort machine; 11  uint version; 12  uint entry; 13  uint phoff; 14  uint shoff; </pre>	<pre> sking@sking:~/Documents/FIFTH SEMESTER IITG/OS Lab/xv6-public\$ objdump -f kernel kernel:      file format elf32-i386 architecture: i386, flags 0x00000112: EXEC_P, HAS_SYMS, D_PAGED start address 0x0010000c </pre>
---	---

Fig 5.5: e\_entry ELF header (left) and entry address of kernel as given by objdump command (right)



## Exercise 6:

- After the experiment, following observations with the reasons were understood:
  - Observation:** All the **8 words** (at 0x00100000) at the point where **BIOS enters the boot-loader** (at address 0x7c00) are **0x00000000**.  
**Reason:** The given address location (i.e 0x00100000) corresponds to an address **1 MB** in the main memory. This is the point **where** the **boot-loader** has to **load the kernel**. As the boot-loader's **first** instruction is yet to be executed **after** the address **0x7c00** (which means that **loading** of kernel into the RAM has **not** started), **hence all** the words are having the value **0** with no meaningful data in them yet.
  - Observation:** At the **second breakpoint** (i.e, at the point where **boot-loader** gives the control to **kernel**), the 8 words now have a **non-zero value**.  
**Reason:** This output is expected since by this time the boot-loader has **loaded** the **kernel** into the memory location starting at address 1 MB in the **RAM** and so the **words at 1 MB** location **contain meaningful** (non-zero) information.

```
(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x00000000 0x00000000 0x00000000 0x00000000
00
0x100010: 0x00000000 0x00000000 0x00000000 0x00000000
00
(gdb) b *0x7d91
Breakpoint 2 at 0x7d91
(gdb) c
Continuing.
The target architecture is set to "i386".
=> 0x7d91: call *0x10018

Thread 1 hit Breakpoint 2, 0x00007d91 in ?? ()
(gdb) x/8x 0x00100000
0x100000: 0x1badb002 0x00000000 0xe4524ffe 0x83e020
0f
0x100010: 0x220f10c8 0x9000b8e0 0x220f0010 0xc0200f
d8
(gdb)
```

Fig 5.6: Output of GDB window for x/8x 0x00100000 command at breakpoints 0x7c00 and 0x7d91.

- Answer to the second question is that **second breakpoint** was set at the instruction corresponding to the address location **0x7d91**. This is the **address of the entry function** (as described in the [Exercise 3](#)) which is **responsible** for giving the **control** to **kernel** and is the **last** instruction to execute of the **boot-loader**.