

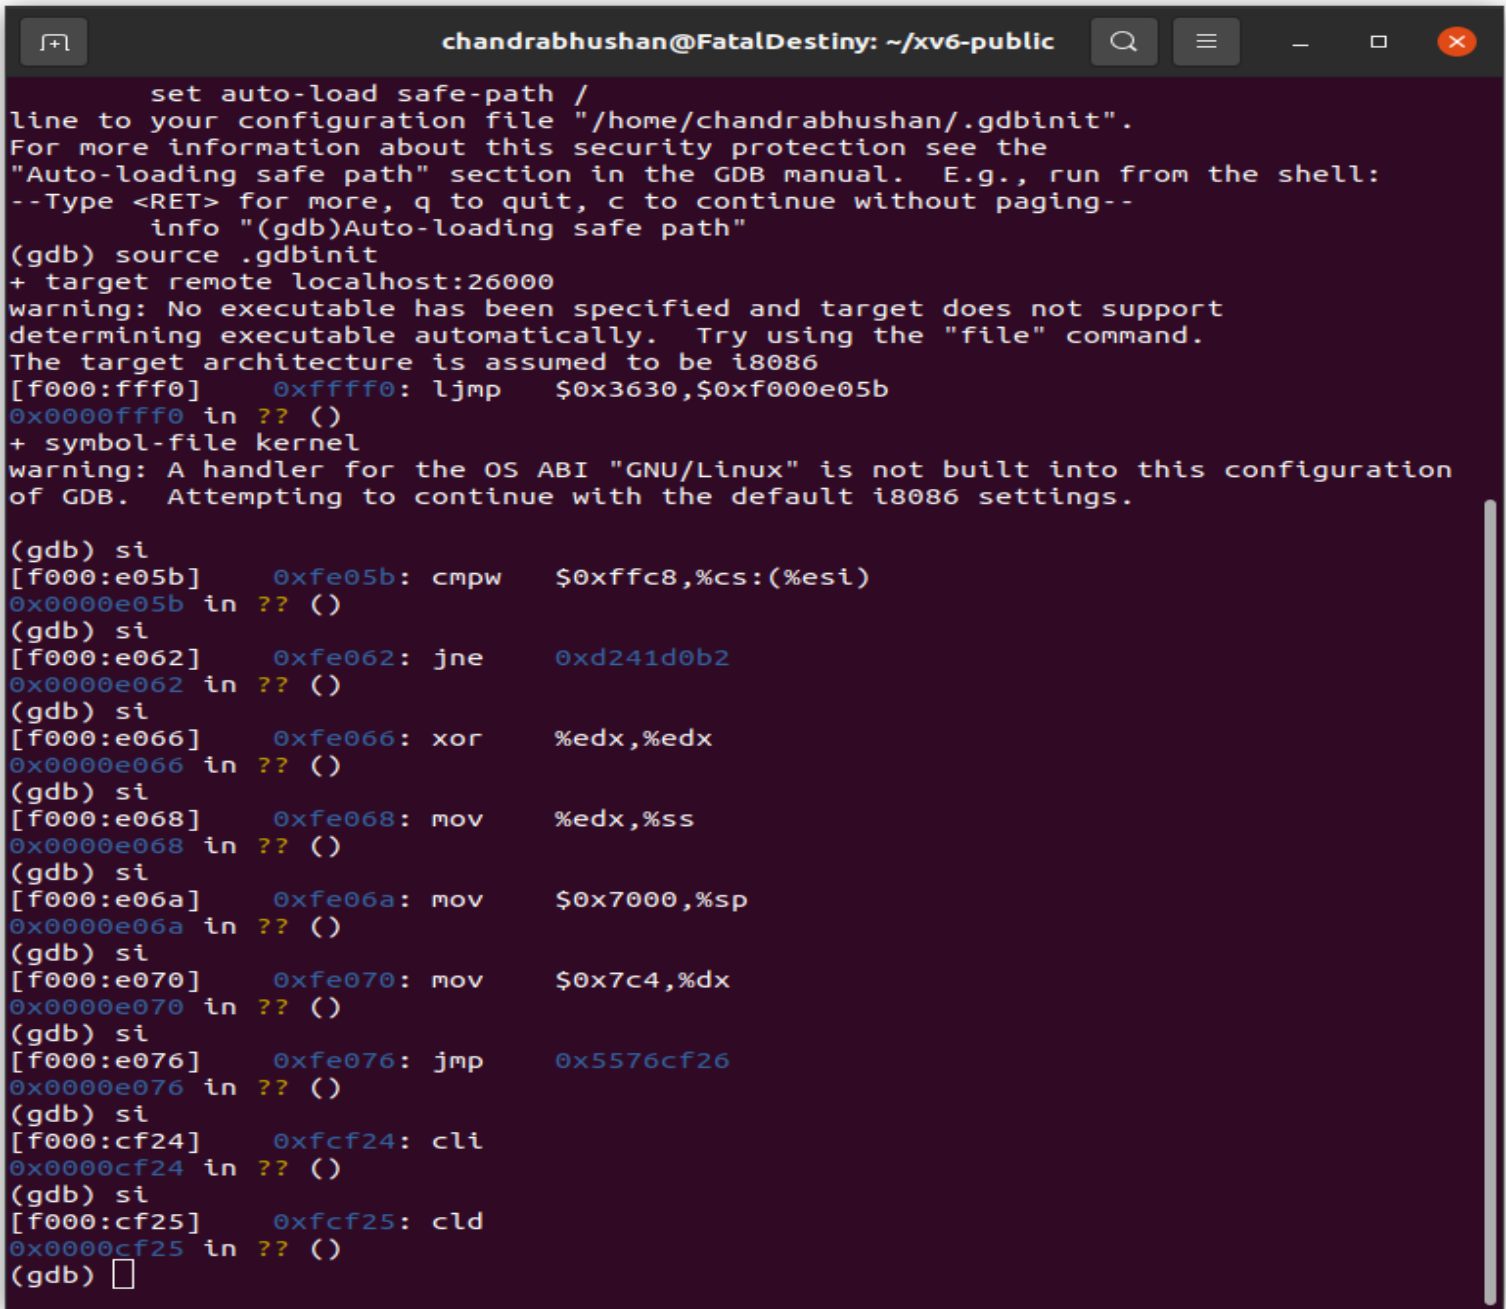
Assignment 1

Name : Chandrabhushan Reddy

Roll Number : 200101027

Exercise 1: Modified code has been submitted in the zip file. Name of the file is exercise1.c

Exercise 2: The following image depicts few instructions of ROM BIOS



Explanation:-

1. [f000:fff0] 0xffff0: ljmp \$0x3630, \$0xf000e05b

- The current location [f000:fff0] = ffff0 is 16 bytes from the end of BIOS ROM region of the memory. Hence, the first thing that the BIOS does is jmp backwards to an earlier location in the BIOS i.e. fe05b

2. [f000:e05b] 0xfe05b: cmpw \$0xffc8, %cs: (%esi)

- Compares ffc8 with the contents of location pointed by code segment and instruction pointer whose values are in cs register and esi register respectively

3. [f000:e062] 0xfe062: jne 0xd241d0b2

- If the previous comparison results in an inequality then jump to location 0xd241d0b2. Else go to the next instruction

4. [f000:e066] 0xfe066: xor %edx, %edx

- Sets the value of edx to zero i.e., %edx is 0

5. [f000:e068] 0xfe068: mov %edx, %ss

- Move contents of edx register to stack segment register (ss). Hence, content of ss register is now 0.

6. [f000:e06a] 0xfe06a: mov \$0x7000, %sp

- Initialised the stack pointer which points to the top of stack to 0x7000

7. [f000:e070] 0xfe070: mov \$0x7c4, %dx

- Sets the value of dx register to 0x7c4

8. [f000:e076] 0xfe076: jmp 0x5576cf26

- Jumps to the location 0x5576cf26

9. [f000:cf24] 0xfcf24: cli

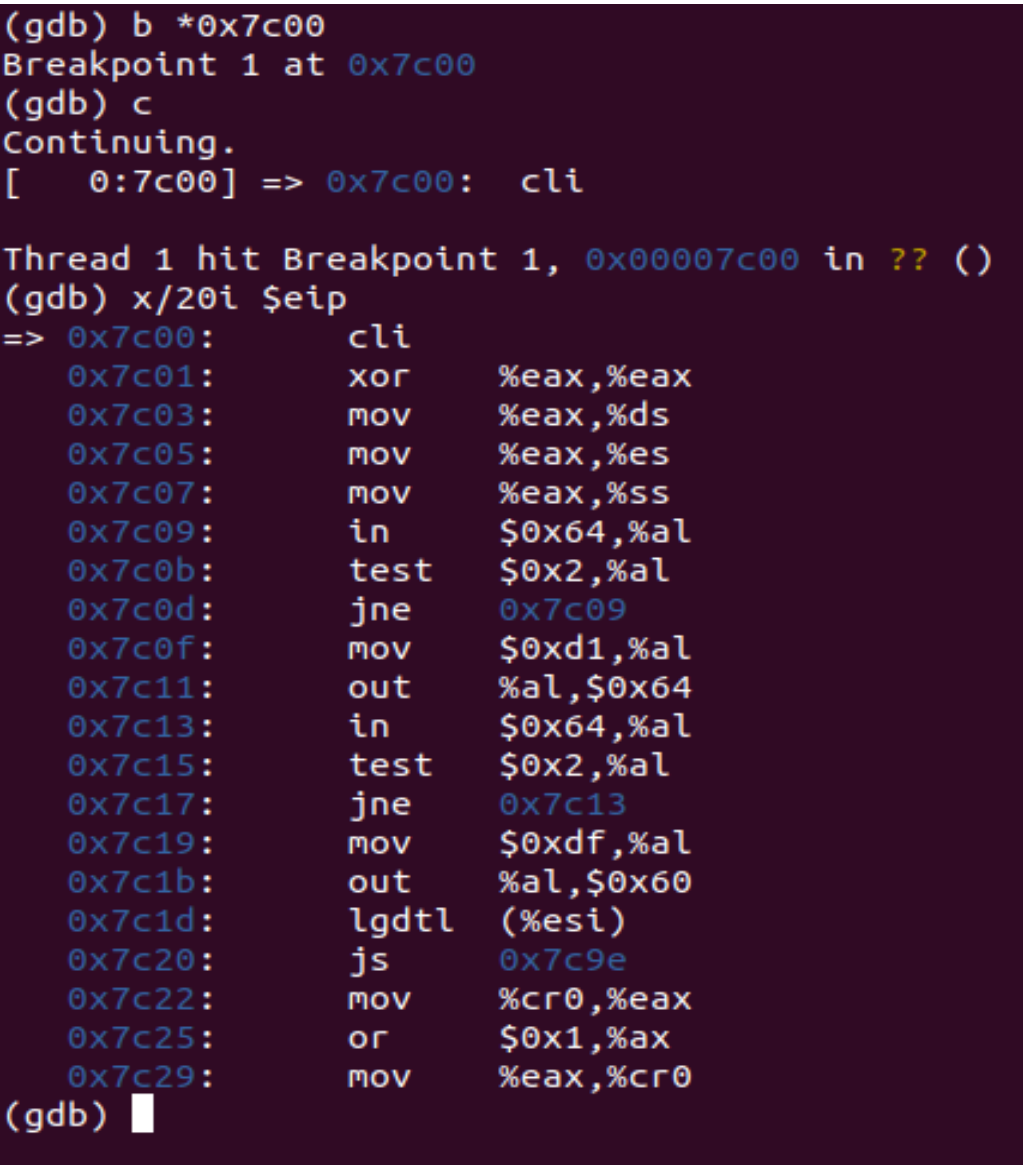
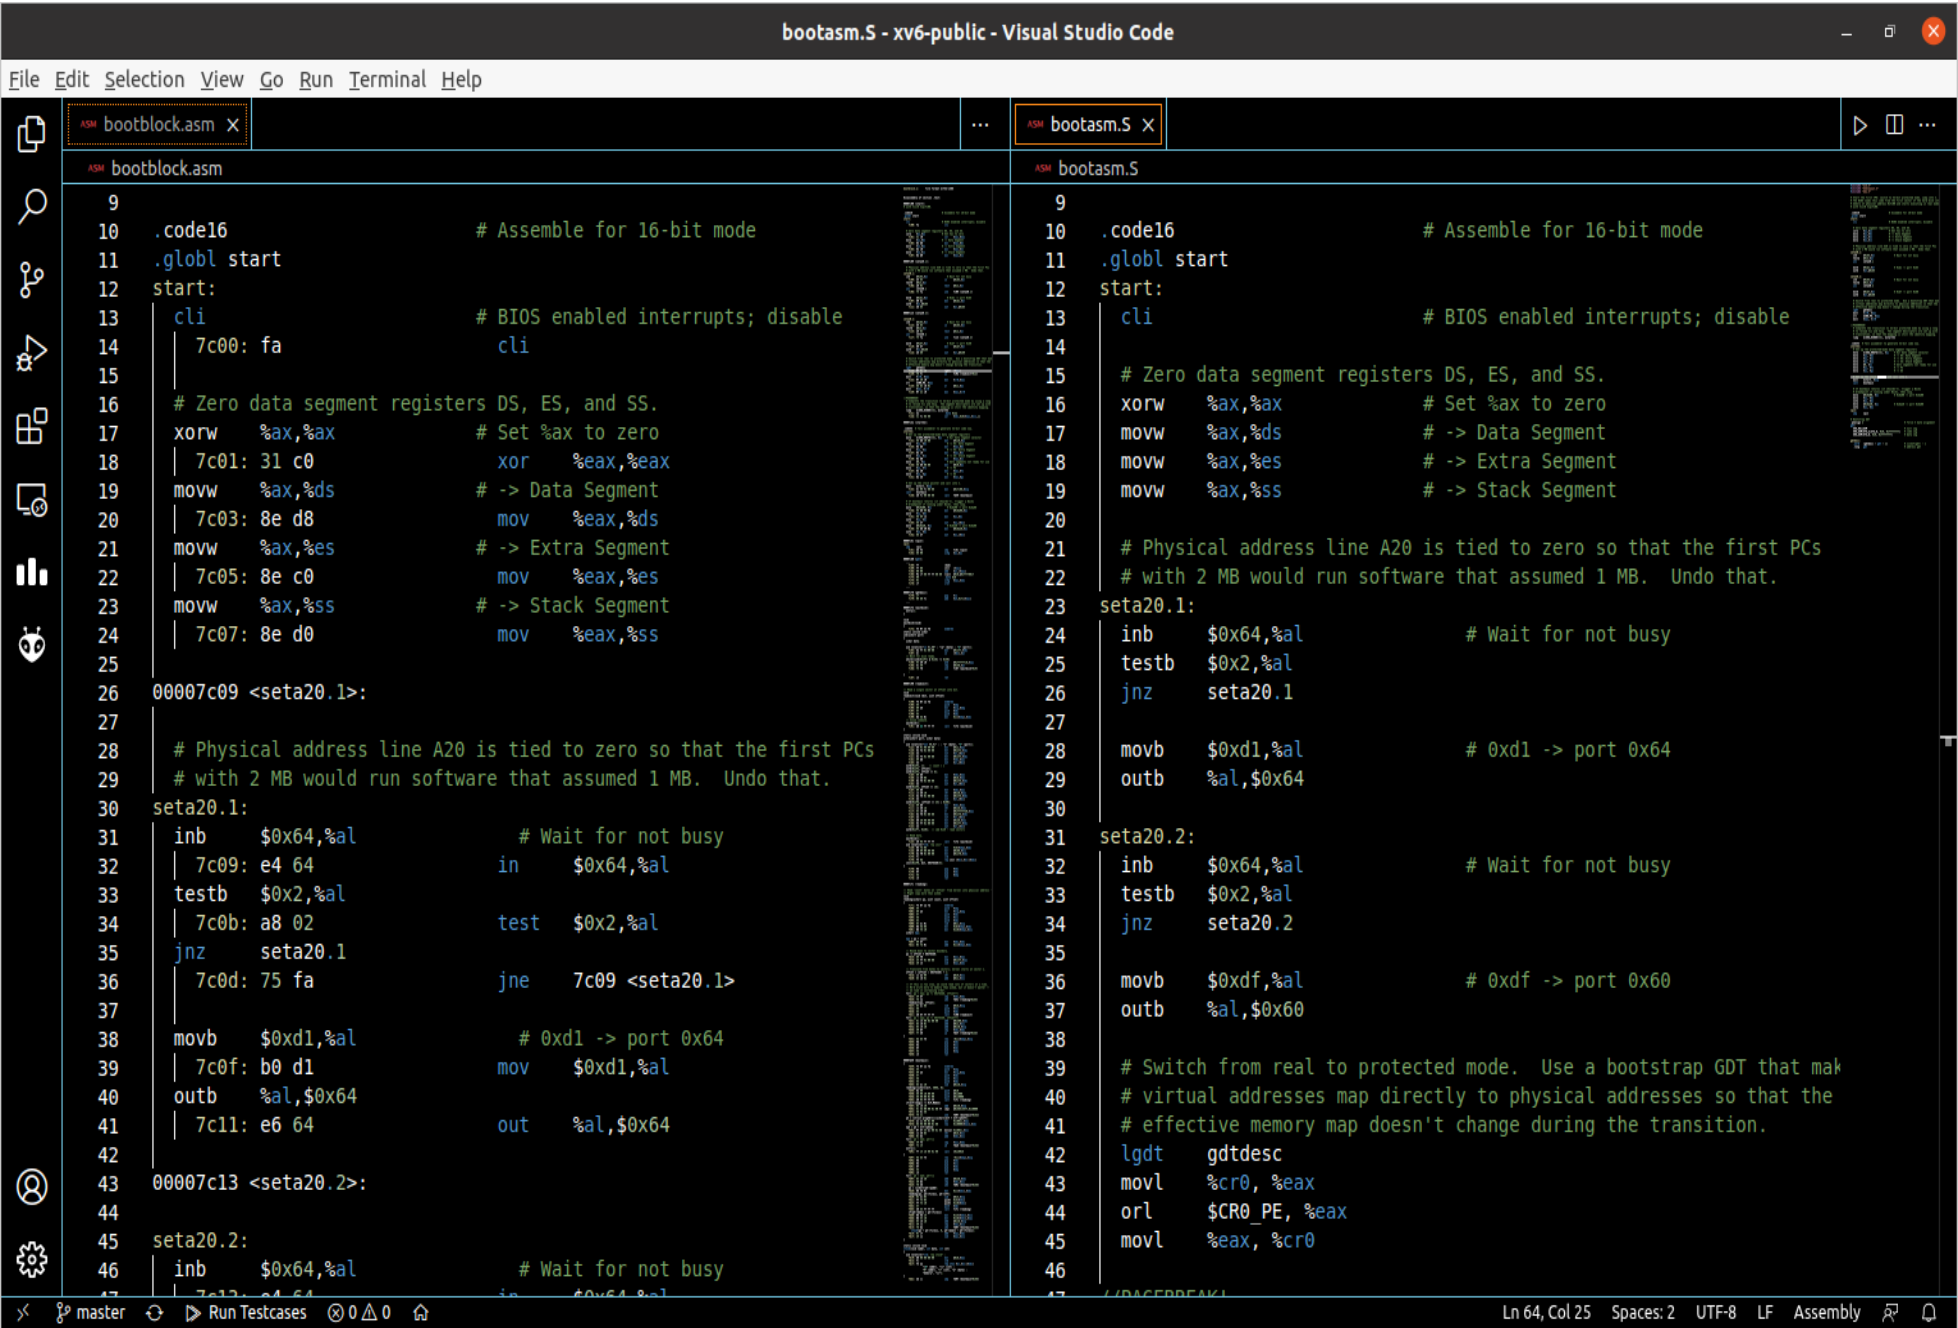
- Clears the interrupt flag. This is done because the CPU should not be interrupted while boot loading.

10. [f000:cf25] 0xfcf25: cld

- Clears the direction flag. This is done for the proper execution of subsequent instructions.

Exercise 3 :

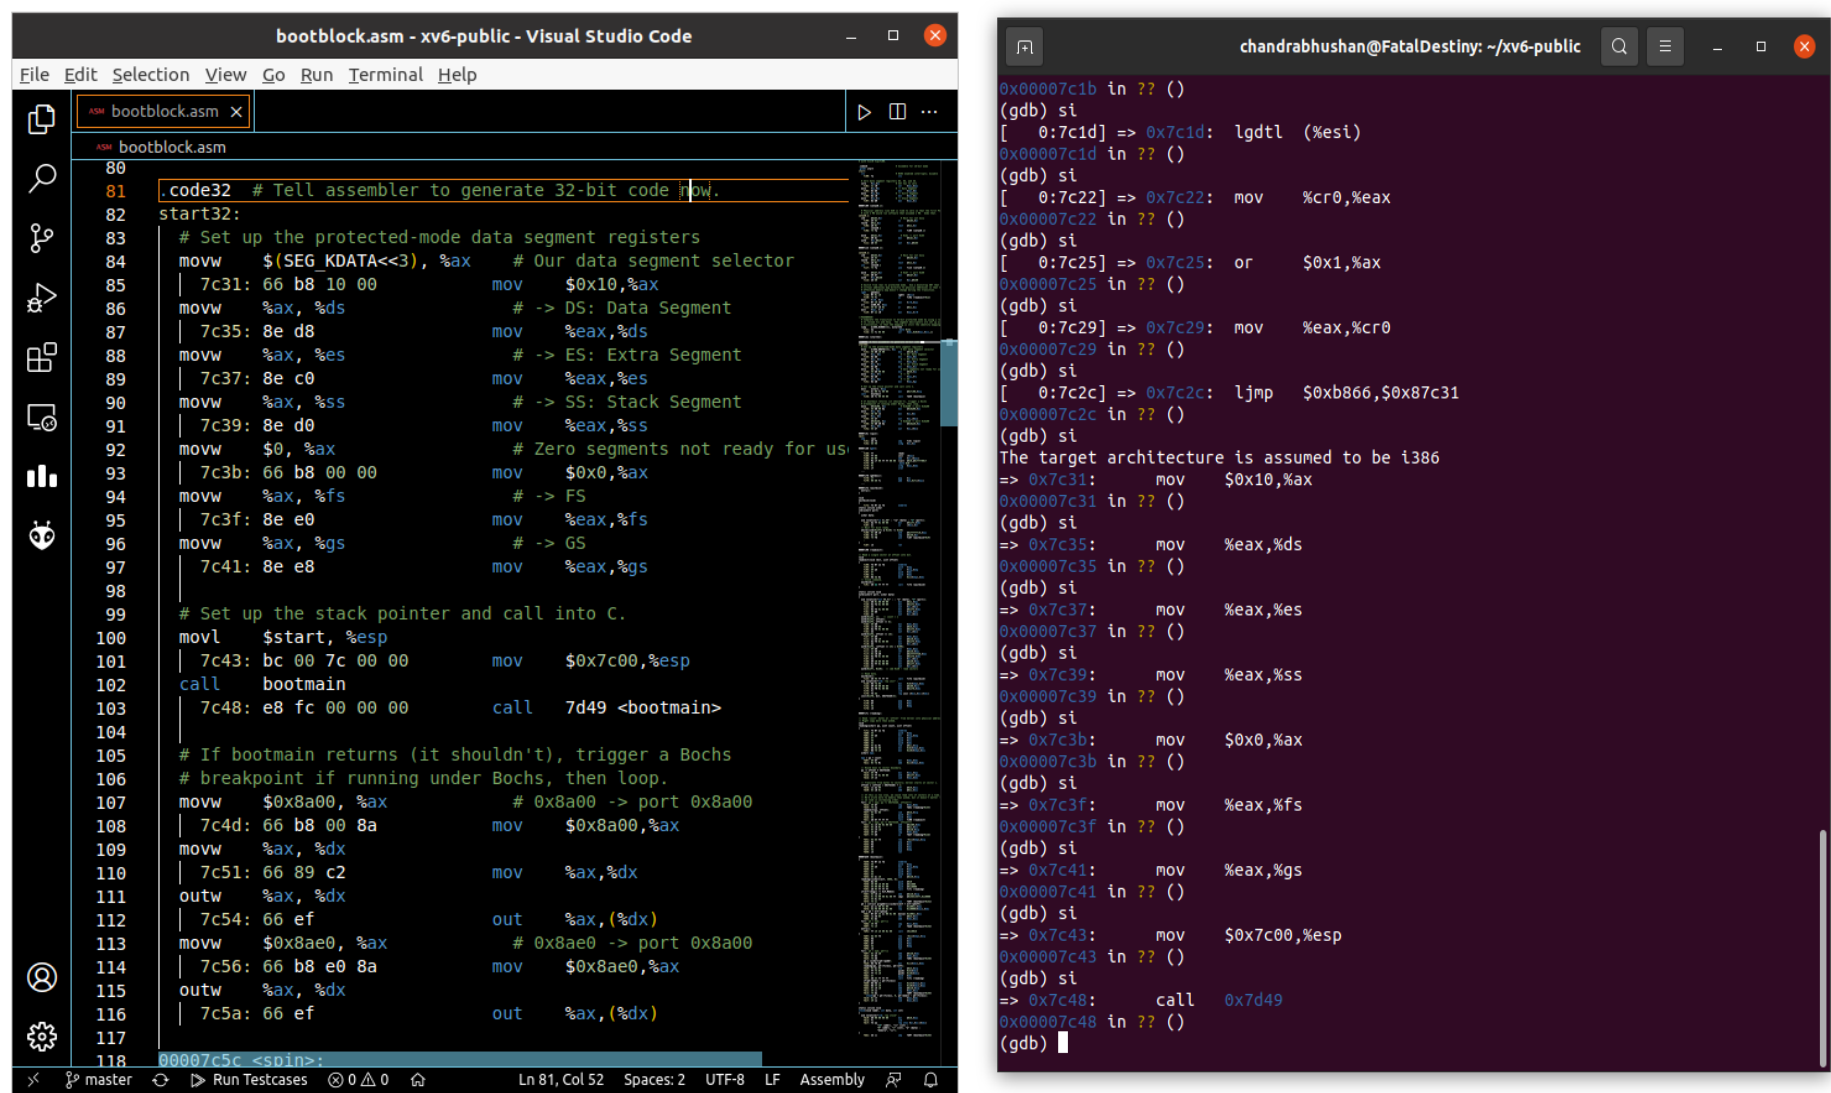
The below figures depict the comparisons between bootasm.S, bootblock.asm and GDB disassembly



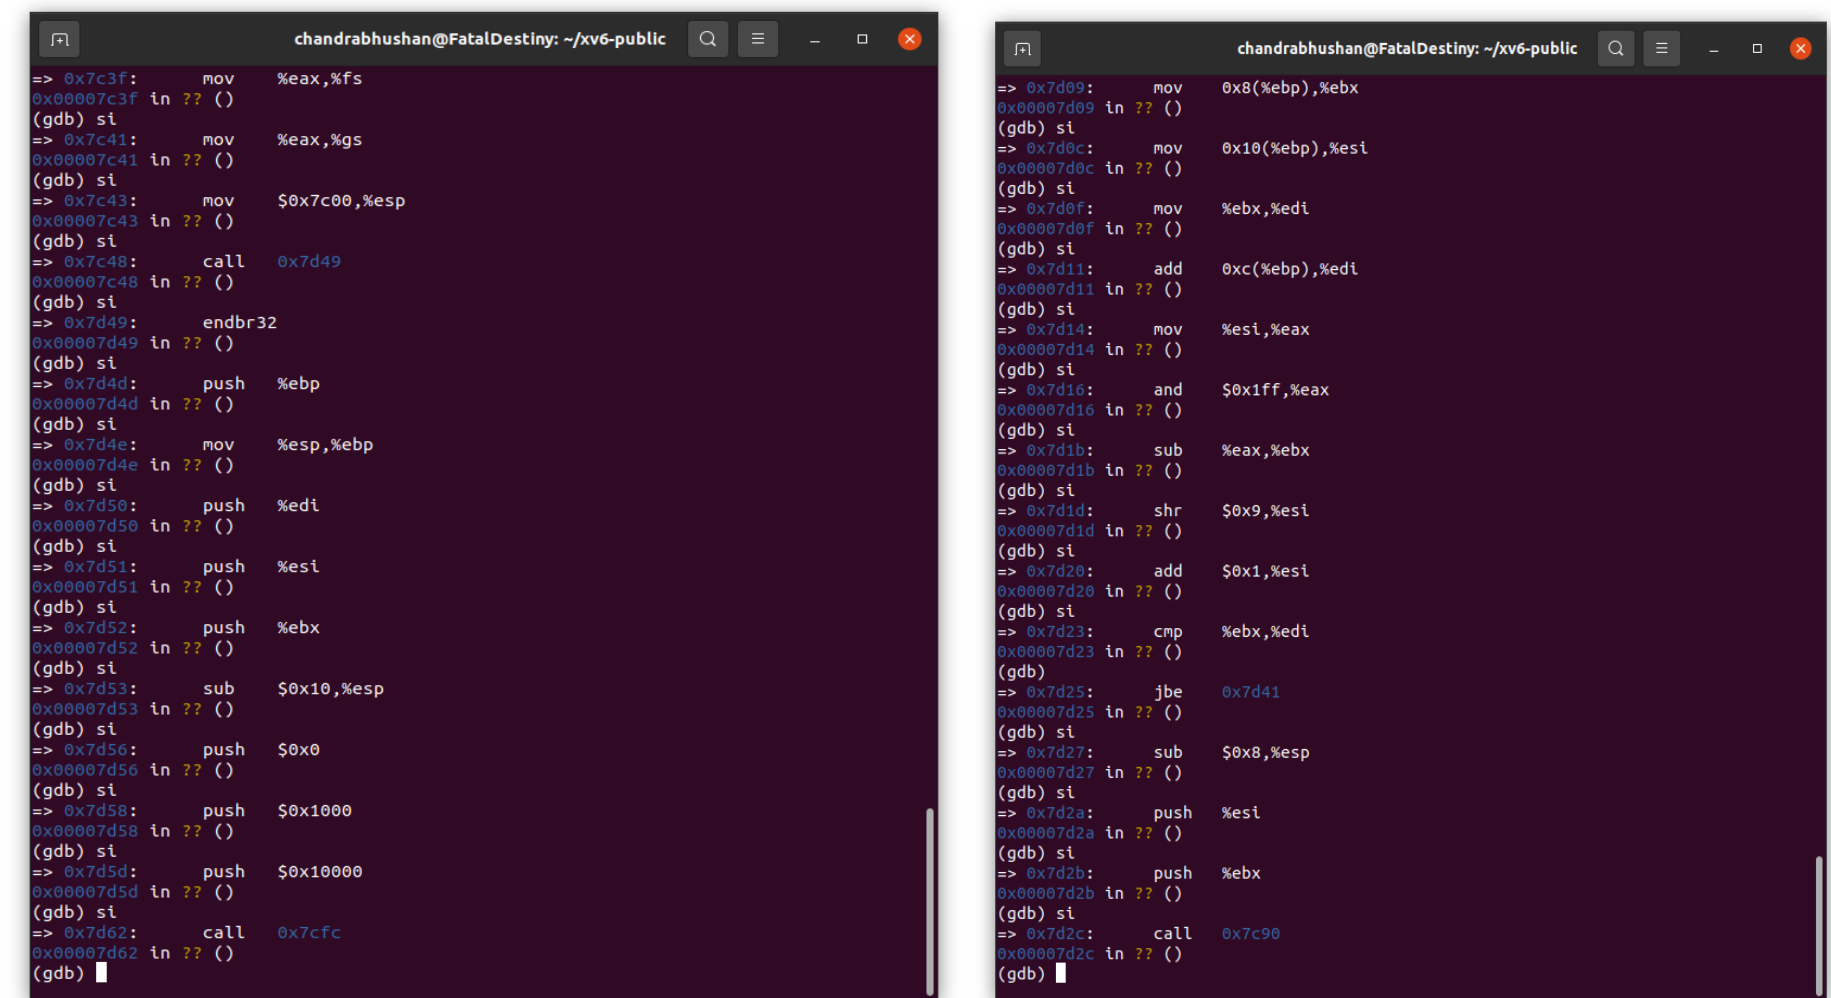
Main differences between bootasm.S and the disassembly on GDB/ bootblock.asm are that:

- **lgdtl gdt desc** instruction in bootasm.S has been expanded to two constituent instructions: **lgdtl (%esi)** and **js 0x7c9e**
- All the instructions starting from cli all the way up to mov %eax, %cr0 are exactly the same except a few changes in the syntax (majorly the suffix of b/w/l standing for byte/word/long word is not present in GDB output)

Code Tracing bootmain() and readsect() in bootmain.c :



In the above bootblock.asm image, line 102 contains the call to bootmain(). This corresponds to the instruction at location 0x7c48 in the GDB call to bootmain (as clearly depicted in the above right image)



The above left image depicts the call to **readseg()** at location **0x7d62** and the right above image depicts the call to **readsect()** at location **0x7d2c** . These images together shows the process of tracing through bootmain() into readseg() and finally into readsect()).


```
bootmain.c - xv6-public - Visual Studio Code
File Edit Selection View Go Run Terminal Help
C bootmain.c x
C bootmain.c > ...
76 // Read 'count' bytes at 'offset' from kernel into physical address
77 // Might copy more than asked.
78 void
79 readseg(uchar* pa, uint count, uint offset)
80 {
81     uchar* epa;
82
83     epa = pa + count;
84
85     // Round down to sector boundary.
86     pa -= offset % SECTSIZE;
87
88     // Translate from bytes to sectors; kernel starts at sector 1.
89     offset = (offset / SECTSIZE) + 1;
90
91     // If this is too slow, we could read lots of sectors at a time.
92     // We'd write more to memory than asked, but it doesn't matter -
93     // we load in increasing order.
94     for(; pa < epa; pa += SECTSIZE, offset++)
95         readsect(pa, offset);
96 }
97
```

```
bootmain.c - xv6-public - Visual Studio Code
File Edit Selection View Go Run Terminal Help
C bootmain.c x
C bootmain.c > ...
39     readseg(pa, ph->filesz, ph->off);
40     if(ph->memsz > ph->filesz)
41         stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
42 }
43
44 // Call the entry point from the ELF header.
45 // Does not return!
46 entry = (void*)(void)(elf->entry);
47 entry();
48 }
49
50 void
51 waitdisk(void)
52 {
53     // Wait for disk ready.
54     while((inb(0x1F7) & 0xC0) != 0x40)
55         ;
56 }
57
58 // Read a single sector at offset into dst.
59 void
60 readsect(void *dst, uint offset)
61 {
62     // Issue command.
63     waitdisk();
64     outb(0x1F2, 1); // count = 1
65     outb(0x1F3, offset);
66     outb(0x1F4, offset >> 8);
67     outb(0x1F5, offset >> 16);
68     outb(0x1F6, (offset >> 24) | 0xE0);
69     outb(0x1F7, 0x20); // cmd 0x20 - read sectors
70
71     // Read data.
72     waitdisk();
73     insl(0x1F0, dst, SECTSIZE/4);
74 }
75
76 // Read 'count' bytes at 'offset' from kernel into physical address
```

The above images show **readseg()** and **readsect()** in bootmain.c at line numbers 79 and 60 respectively.

Assembly level instructions corresponding to each instruction in readsect():

```
Line 63 : => 0x7c9c:    call    0x7c7e
0x000007c9c in ?? ()
```

```
=> 0x7cbf:    mov     %ebx,%eax
0x000007cbf in ?? ()
(gdb) si
=> 0x7cc1:    shr     $0x10,%eax
0x000007cc1 in ?? ()
(gdb) si
=> 0x7cc4:    mov     $0x1f5,%edx
0x000007cc4 in ?? ()
(gdb) si
=> 0x7cc9:    out     %al,(%dx)
0x000007cc9 in ?? ()
```

```
Line 64 : 0x7ca1:    mov     $0x1,%eax
0x7ca6:    mov     $0x1f2,%edx
(gdb) si
=> 0x7c7e:    endbr32
```

```
(gdb) si
=> 0x7cca:    mov     %ebx,%eax
0x000007cca in ?? ()
(gdb) si
=> 0x7ccc:    shr     $0x18,%eax
0x000007ccc in ?? ()
(gdb) si
=> 0x7ccf:    or      $0xfffffffffe0,%eax
0x000007ccf in ?? ()
(gdb) si
=> 0x7cd2:    mov     $0x1f6,%edx
0x000007cd2 in ?? ()
(gdb) si
=> 0x7cd7:    out     %al,(%dx)
0x000007cd7 in ?? ()
```

```
Line 65 : => 0x7cac:    mov     $0x1f3,%edx
0x000007cac in ?? ()
(gdb) si
=> 0x7cb1:    mov     %ebx,%eax
0x000007cb1 in ?? ()
(gdb) si
=> 0x7cb3:    out     %al,(%dx)
0x000007cb3 in ?? ()
```

```
(gdb) si
=> 0x7cd8:    mov     $0x20,%eax
0x000007cd8 in ?? ()
(gdb) si
=> 0x7cdd:    mov     $0x1f7,%edx
0x000007cdd in ?? ()
(gdb) si
=> 0x7ce2:    out     %al,(%dx)
0x000007ce2 in ?? ()
```

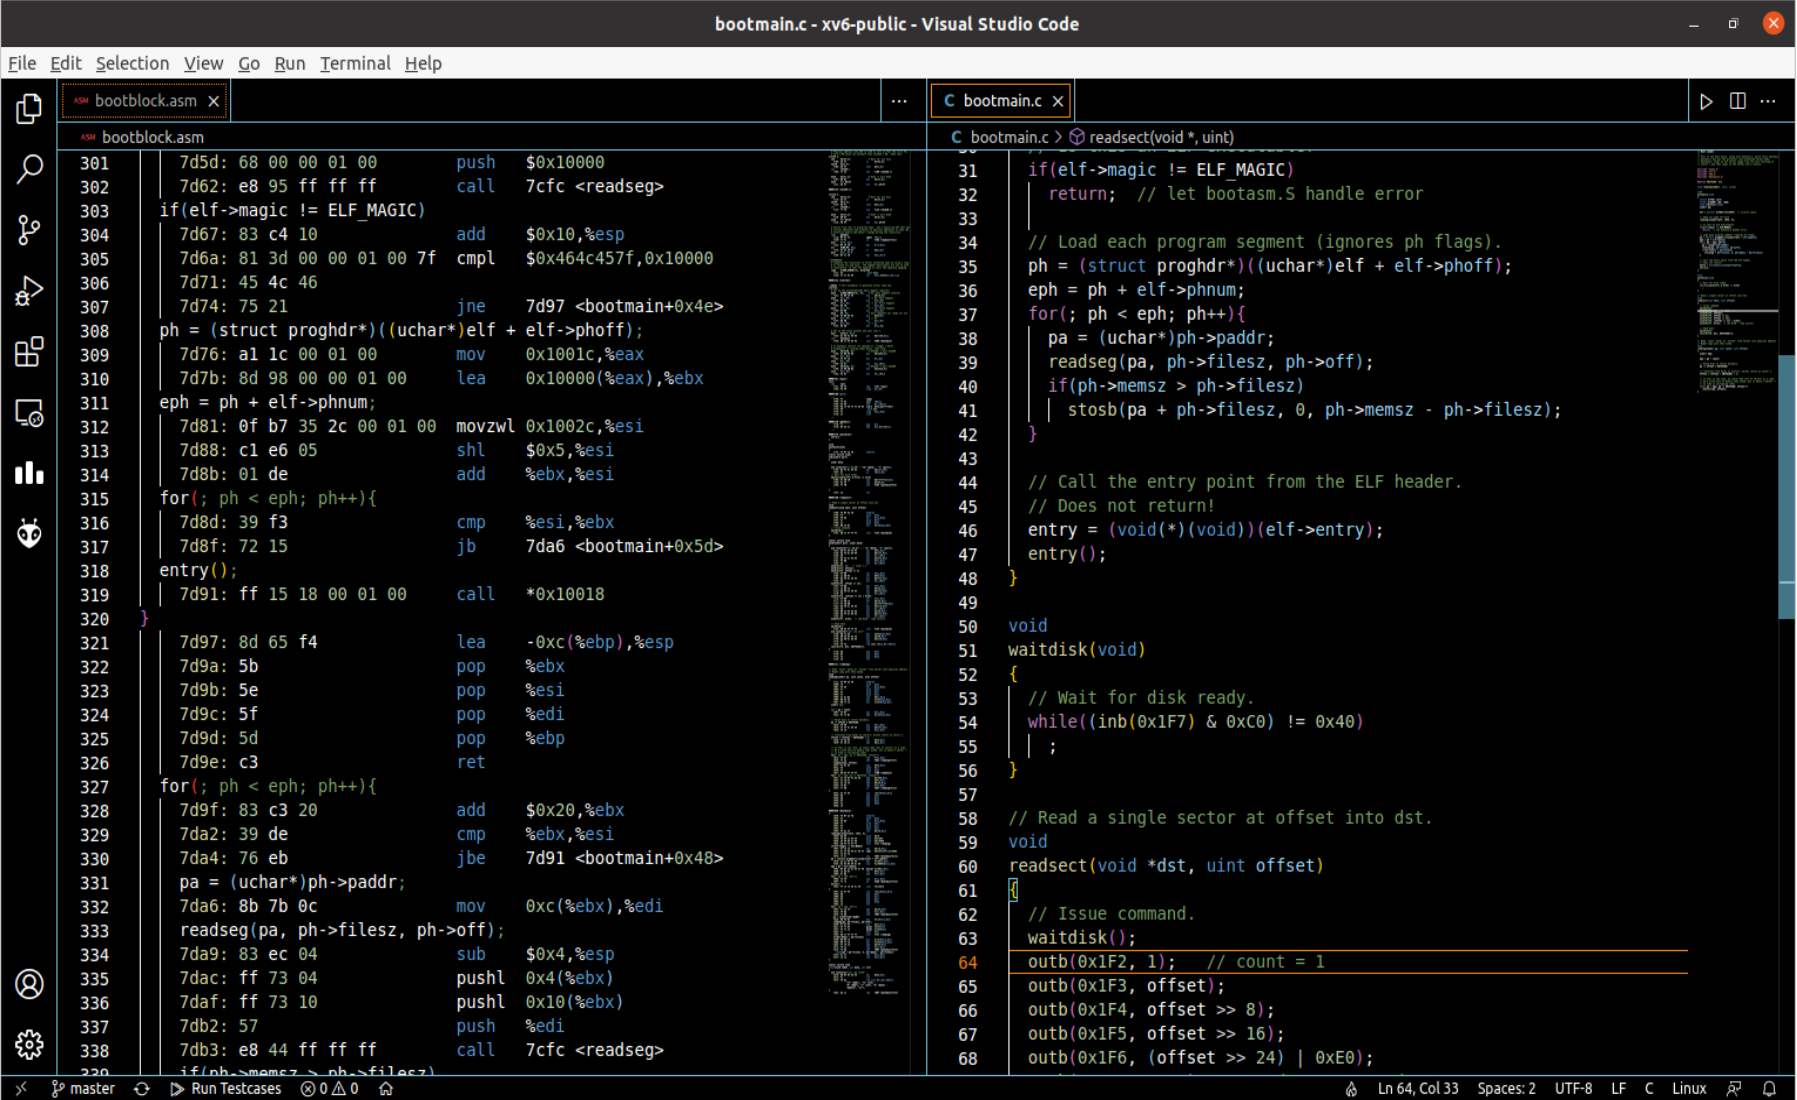
```
Line 66 : => 0x7cb4:    mov     %ebx,%eax
0x000007cb4 in ?? ()
(gdb) si
=> 0x7cb6:    shr     $0x8,%eax
0x000007cb6 in ?? ()
(gdb) si
=> 0x7cb9:    mov     $0x1f4,%edx
0x000007cb9 in ?? ()
(gdb) si
=> 0x7cbe:    out     %al,(%dx)
0x000007cbe in ?? ()
```

```
Line 72 : (gdb) si
=> 0x7ce3:    call    0x7c7e
0x000007ce3 in ?? ()
```

```
(gdb) si
=> 0x7ce8:      mov     0x8(%ebp),%edi
0x000007ce8 in ?? ()
(gdb) si
=> 0x7ceb:      mov     $0x80,%ecx
0x000007ceb in ?? ()
(gdb) si
=> 0x7cf0:      mov     $0x1f0,%edx
0x000007cf0 in ?? ()
(gdb) si
=> 0x7cf5:      cld
0x000007cf5 in ?? ()
(gdb) si
=> 0x7cf6:      rep insl (%dx),%es:(%edi)
0x000007cf6 in ?? ()
```

Line 73 :

Details about for loop that reads remaining sectors of kernel into memory :



The above image shows the for loop that reads remaining sectors of the kernel into the memory. The left image shows the assembly code and the right image shows the source code.

The code on line 46-47 (in right side source file) whose assembly is on line 318-319 (left side assembly file) of above image would be executed once the for loop is over. This line is the final line in execution of boot loader and after this the control is given to the OS.

```
(gdb) b *0x7d91
Breakpoint 2 at 0x7d91
(gdb) c
Continuing.
=> 0x7d91:      call    *0x10018

Thread 1 hit Breakpoint 2, 0x000007d91 in ?? ()
(gdb) si
=> 0x10000c:    mov     %cr4,%eax
0x0010000c in ?? ()
(gdb) si
=> 0x10000f:    or      $0x10,%eax
0x0010000f in ?? ()
(gdb) si
=> 0x100012:    mov     %eax,%cr4
0x00100012 in ?? ()
(gdb) si
=> 0x100015:    mov     $0x109000,%eax
0x00100015 in ?? ()
(gdb) █
```

The breakpoint is set at 0x7d91 (line 319) as shown in the above image

Question 1 :

Following image shows the last lines executed before switching to the protected mode. The last line to be executed is line 51. Before that the boot loader deactivates the line A20 (to access 2MB of RAM instead of default 1MB) and also initializes the descriptor table which would be helpful in accessing the virtual memory in the protected mode. After it's done then the line 51 causes the control to jump to the protected mode.

```
|
| # Switch from real to protected mode. Use a bootstrap GDT that makes
| # virtual addresses map directly to physical addresses so that the
| # effective memory map doesn't change during the transition.
| lgdt    gdt_desc
| 7c1d: 0f 01 16          lgdtl   (%esi)
| 7c20: 78 7c          js      7c9e <readsect+0xe>
| movl    %cr0, %eax
| 7c22: 0f 20 c0          mov     %cr0, %eax
| orl     $CR0_PE, %eax
| 7c25: 66 83 c8 01       or      $0x1, %ax
| movl    %eax, %cr0
| 7c29: 0f 22 c0          mov     %eax, %cr0
|
//PAGEBREAK!
| # Complete the transition to 32-bit protected mode by using a long jmp
| # to reload %cs and %eip. The segment descriptors are set up with no
| # translation, so that the mapping is still the identity mapping.
| ljmp    $(SEG_KCODE<<3), $start32
| 7c2c: ea          .byte 0xea
| 7c2d: 31 7c 08 00       xor     %edi, 0x0(%eax, %ecx, 1)
```

Assembly code to switch from 16 bit Real mode to 32 bit Protected mode.

Question 2 :

From the above discussion it is clear that the last instruction of boot loader would be to call the entry function (on line 47 in bootmain.c). Thus a breakpoint is set for the corresponding assembly instruction and the command si is used to trace the next instruction (first instruction of the kernel) and following conclusions are made based on the screenshot below:

Last instruction of boot loader: call *0x10018

First instruction of kernel: mov %cr4, %eax

```
(gdb) b *0x7d91
Breakpoint 2 at 0x7d91
(gdb) c
Continuing.
=> 0x7d91:      call    *0x10018

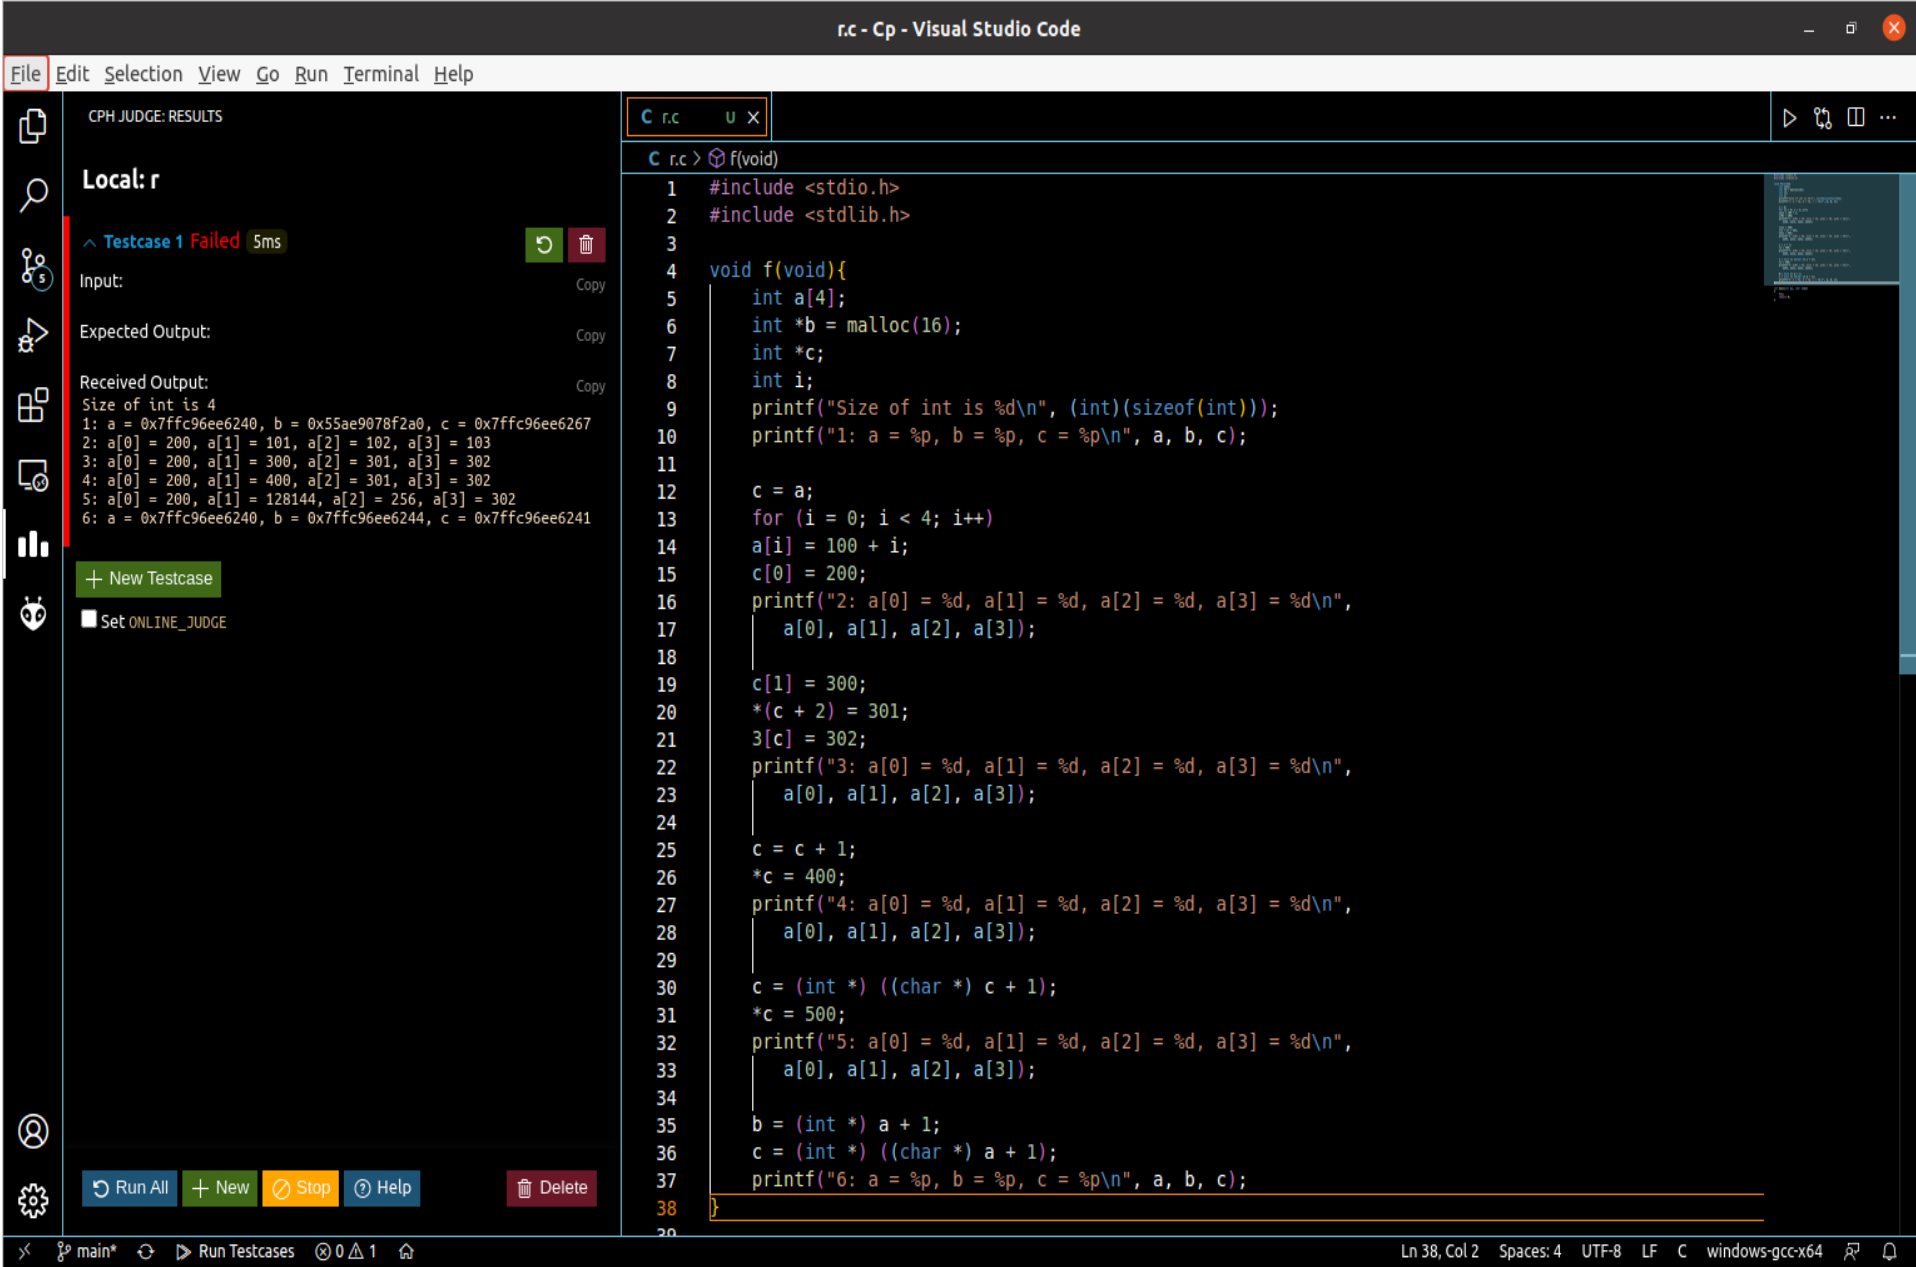
Thread 1 hit Breakpoint 2, 0x00007d91 in ?? ()
(gdb) si
=> 0x10000c:    mov     %cr4,%eax
0x0010000c in ?? ()
```

Question 3 :

It is clear from the following image that the boot loader runs a for loop starting at ph and ending at eph – 1. The value of ph is obtained from the ELF header and the number of iterations are given by elf->phnum, thus eph = ph + elf->phnum.

```
// Load each program segment (ignores ph flags).
ph = (struct proghdr*)((uchar*)elf + elf->phoff);
eph = ph + elf->phnum;
for(; ph < eph; ph++){
    pa = (uchar*)ph->paddr;
    readseg(pa, ph->filesz, ph->off);
    if(ph->memsz > ph->filesz)
        stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
}
```

Exercise 4 :



Below is the line by line explanation of each line of output for above code :

Explanation for 1st output line: The variables a and c belong to stack memory address and thus store close addresses. Malloc is used to dynamically allocate address in the heap space and so the value stored in b is having a totally different address.

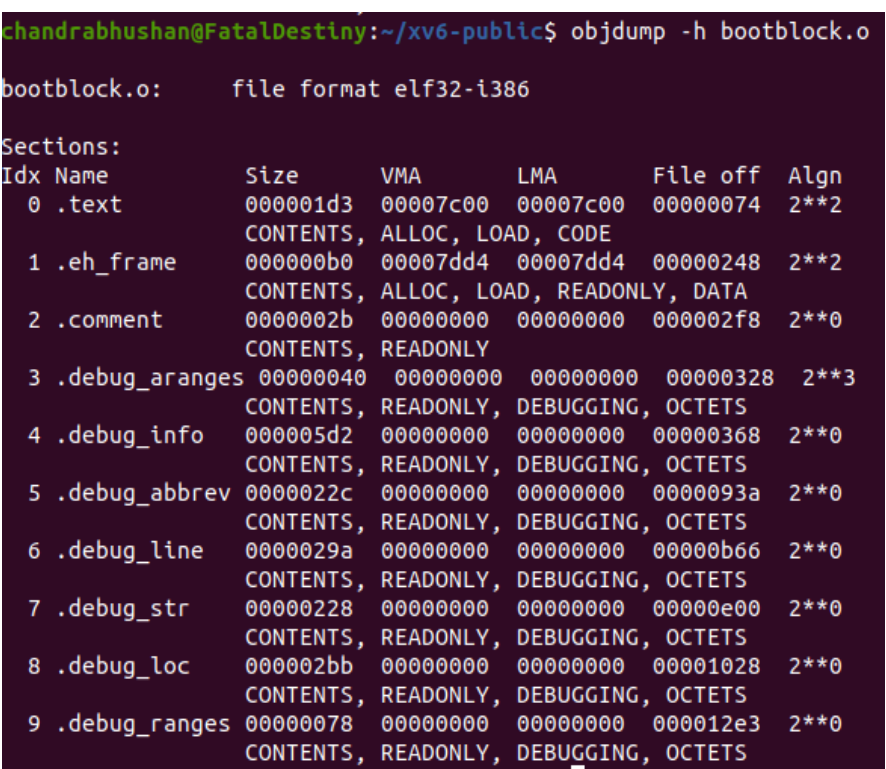
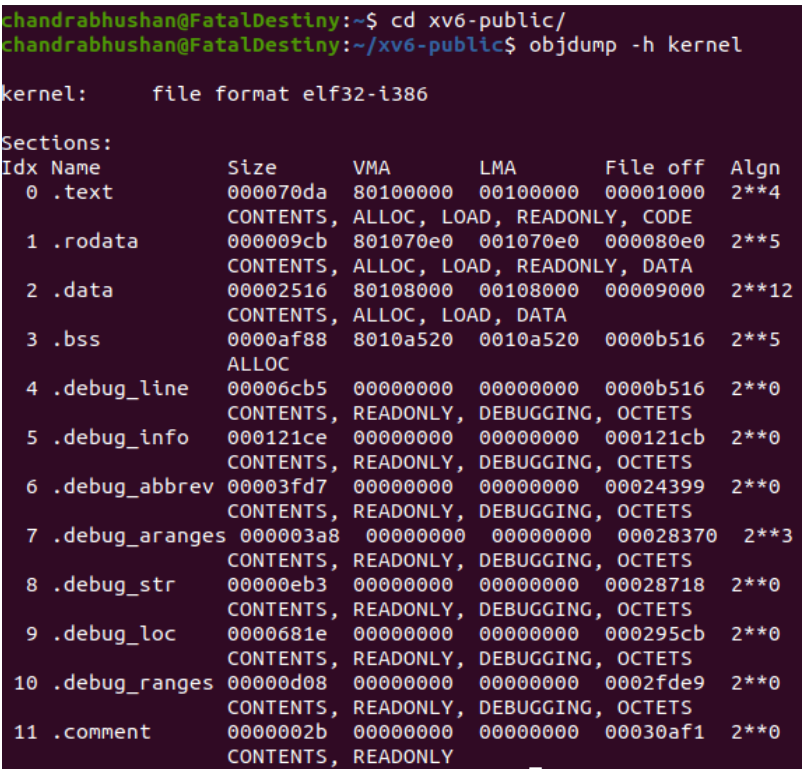
Explanation for 2nd output line: After line 13, the pointer c stores the address of array a. Thus line 15 changes the value of previously assigned a[0] of 100 to 200 and so in the output the value of a[0] is 200 whereas the other values of the array are same as before.

Explanation for 3rd output line: Line 20 changes the value of a[1] to 300 (as c is still storing the address of array a). Lines 21 and 22 are just two other ways of dereferencing the elements of an array and change a[2] and a[3] to 301 and 302 respectively. Thus the only difference between output of line 2 and 3 is the values of a[1], a[2] and a[3].

Explanation for 4th output line: Line 26 increases the address in c by 4 bytes (since c is of the type int*). Thus c now points to address of a[1] and Line 27 changes a[1] to 400. Hence the output in line 4 is same as that of line 2 except the value of a[1] is now 400.

Explanation for 5th output line: Line 31 increases the address in c by 1 byte (since c is changed to char* before incrementing). Now c doesn't point to any specific element in a and thus the value of 500 is assigned partially between a[1] and a[2] and it changes both the values in a corrupted way.

Explanation for 6th output line: Line 36 stores the address of a[1] in b and line 37 stores the address of a incremented by 1 byte (similar to line 31). Thus the value of a is same as output of line 1 and value of b is a + 4 bytes and value of c is a + 1 byte.



Above figures show various program sections of the kernel (left above image) and bootblock.o (right above image) binaries. The output contains the following columns:

- **Name:** Name of the program section. Eg: .text contains program instructions, .data contains initialized global variables, etc.
- **Size:** Size of the program section in bytes.
- **VMA:** Link address of the program section. This is the address at which the program expects to be executed.

- **LMA:** Load address of program section. This is the address into which the program section is actually loaded.
- **File off:** Offset of the section from beginning of file in disk.
- **Align:** Alignment to accommodate various data types.

Exercise 6 :

Below is the hypothesis made by observations:

When BIOS just enters the bootloader no useful information is present at 0x100000. However, from Exercise 3 Question 3 we know that the boot loader stores the kernel (ELF Header) starting from address 0x100000. Hence, when the bootloader enter the kernel, the contents at 0x100000 contains the kernel image.

We can confirm this using GDB as shown in the following figure. We set two breakpoints- one at the entry of bootloader and another at the entry of kernel. At both breakpoints we view 8 words at location 0x100000. From the GDB tracing we can conclude:

1. When the BIOS enters bootloader, the contents of 8 words at 0x100000 is all zero
2. Once the bootloader is executed and enters kernel, the contents of 8 words at 0x100000 contain different information.
3. This is because the bootloader stores the ELF header at 0x100000 as seen in the figure pasted in exercise 3.

```
chandrabhushan@FatalDestiny: ~/xv6-public
warning: File "/home/chandrabhushan/xv6-public/.gdbinit" auto-loading has been declined by your `auto-load safe-path' set to "$debugdir:$datadir/auto-load".
To enable execution of this file add
    add-auto-load-safe-path /home/chandrabhushan/xv6-public/.gdbinit
line to your configuration file "/home/chandrabhushan/.gdbinit".
To completely disable this security protection add
    set auto-load safe-path /
line to your configuration file "/home/chandrabhushan/.gdbinit".
For more information about this security protection see the
"Auto-loading safe path" section in the GDB manual.  E.g., run from the shell:
--Type <RET> for more, q to quit, c to continue without paging--
info "(gdb)Auto-loading safe path"
(gdb) source .gdbinit
+ target remote localhost:26000
warning: No executable has been specified and target does not support
determining executable automatically.  Try using the "file" command.
The target architecture is assumed to be i8086
[f000:fff0] 0xffff0: jmp $0x3630,$0xf000e05b
0x0000fff0 in ?? ()
+ symbol-file kernel
warning: A handler for the OS ABI "GNU/Linux" is not built into this configuration
of GDB.  Attempting to continue with the default i8086 settings.

(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) b *0x10000c
Breakpoint 2 at 0x10000c
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Thread 1 hit Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/8x 0x100000
0x100000: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0x100010: 0x00000000 0x00000000 0x00000000 0x00000000
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x10000c: mov %cr4,%eax

Thread 1 hit Breakpoint 2, 0x0010000c in ?? ()
(gdb) x/8x 0x100000
0x100000: 0x1badb002 0x00000000 0xe4524ffe 0x83e0200f
0x100010: 0x220f10c8 0x9000b8e0 0x220f0010 0xc0200fd8
(gdb) █
```